Training

1. **Data Loading and Preprocessing:**
   - Data and labels are loaded from multiple .npy files.
   - Features are normalized to the range [0, 1] to ensure stable training.
   - Data is reshaped to include a single channel dimension, as the images are grayscale.

2. **Target Binning:**
   - Continuous labels are discretized into a specified number of bins (default: 50).
   - Each bin represents a range of values, transforming the regression problem into classification.

3. **Model Architecture:**
   A CNN model is constructed using TensorFlow's Keras API. Key design elements include:
   - Three Convolutional Blocks:
     - Sequential convolutional layers with increasing filter sizes (64 → 128 → 256).
     - Batch normalization for stabilizing training.
     - Activation functions (ReLU) for non-linearity.
     - Max-pooling for spatial down-sampling and dropout for regularization.
   - Fully Connected Layers:
     - Dense layers with L2 regularization to prevent overfitting.
     - Batch normalization and dropout for further stability.
     - Final softmax activation for multi-class classification.
   - The model is compiled with the Adam optimizer, categorical cross-entropy loss, and accuracy metric.

4. **Training:**
   - The dataset is split into training and testing sets, with further division of the training set for validation.
   - Early stopping and learning rate reduction callbacks are used to prevent overfitting and adapt learning rates during training.
   - The model is trained over a maximum of 100 epochs with a batch size of 32.

5. **Evaluation:**
   - The training history is plotted, showing loss and accuracy trends for both training and validation sets.
   - Model weights and bin edges are saved for future inference.

Results:

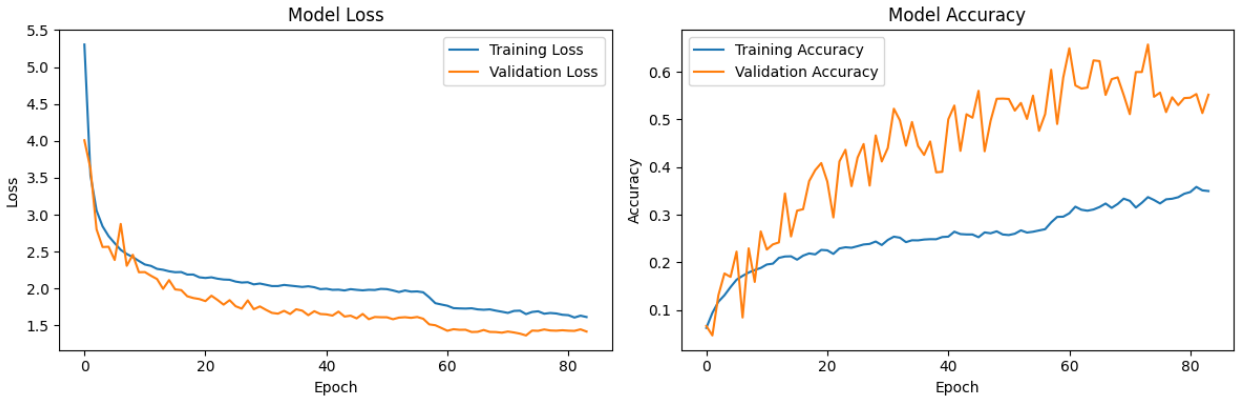

Inference

1. **Model Reconstruction:**
    - The model architecture is recreated to match the training configuration.
    - Saved weights and bin edges are loaded to prepare the model for inference.
2. **Prediction:**
    - Images are preprocessed similarly to the training pipeline (normalization and reshaping).
    - Predictions are made on the test data, and the output class indices are mapped back to the continuous scale using the loaded bin edges.
3. **Evaluation Metrics:**
    - Mean Absolute Error (MAE): Measures average prediction error.
    - Mean Squared Error (MSE): Penalizes larger errors more heavily than MAE.
    - Accuracy Within Thresholds: The percentage of predictions falling within ±0.5, ±1.0, and ±2.0 of the true values.
4. **Visualization:**
    - Random samples of predictions are visualized alongside the true values.
    - A scatter plot illustrates the relationship between predicted and true values, with a diagonal reference line for ideal predictions.

Results:

Test Results:
Mean Absolute Error: 0.6393
Mean Squared Error: 0.6901
Accuracy (within ±0.5): 49.12%
Accuracy (within ±1.0): 75.95%
Accuracy (within ±2.0): 98.55%

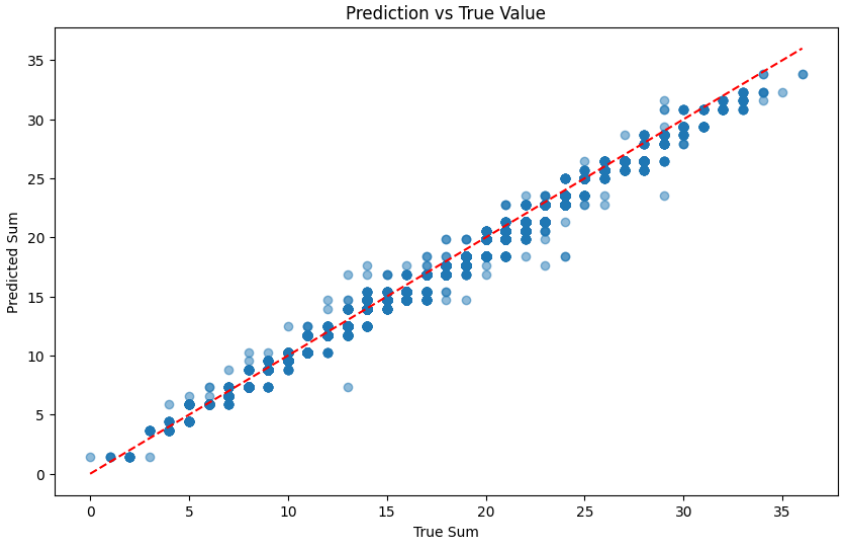

Demo results: